

Work Note of RDMA Study

Zhuangdi Zhu
zhuangdizhu@yahoo.com
2015-01-27

Target: Enable to virtualize InfiniBand adapters on containers

Progress:

Step1: install EPEL

```
[root@zhuangdizhu2]# wget http://dl.fedoraproject.org/pub/epel/6/x86_64/epel-release-6-8.noarch.rpm
```

```
[root@zhuangdizhu2]# sudo rpm -Uvh epel-release-6*.rpm
```

Step2: install docker

```
[root@zhuangdizhu2 init.d]# yum -y install docker-io
```

```
[root@zhuangdizhu2 init.d]# cat /var/log/docker
```

docker - version 1.5

Copyright 2003, Ben Jansens <ben@orodu.net>

Usage: docker [OPTIONS]

Options:

-help Show this help.

-display DISPLAY The X display to connect to.

-border The width of the border to put around the
 system tray icons. Defaults to 1.

-vertical Line up the icons vertically. Defaults to
 horizontally.

-wmaker WindowMaker mode. This makes docker a
 fixed size (64x64) to appear nicely in
 in WindowMaker.

Note: In this mode, you have a fixed
 number of icons that docker can hold.

-iconsize SIZE The size (width and height) to display
 icons as in the system tray. Defaults to
 24.

\nTue Jan 27 06:21:27 CST 2015\n

2015/01/27 06:21:27 docker daemon: 1.3.2 39fa2fa/1.3.2; execdriver: native; graphdriver:

[74c93085] +job serveapi(unix:///var/run/docker.sock)

[info] WARNING: You are running linux kernel version 2.6.32-431.el6.x86_64, which might be unstable running docker.
Please upgrade your kernel to 3.8.0.

[info] Listening for HTTP on unix (/var/run/docker.sock)

[74c93085] +job init_networkdriver()

[74c93085] -job init_networkdriver() = OK (0)

```
[info] Loading containers:
```

```
[info] : done.
```

```
[74c93085] +job acceptconnections()
```

```
[74c93085] -job acceptconnections() = OK (0)
```

```
[root@zhuangdizhu2 init.d]#
```

Step3: start docker

```
[root@zhuangdizhu2 init.d]# cd /etc/init.d
```

```
[root@zhuangdizhu2 init.d]# ./docker start
```

```
Starting docker: [ OK ]
```

```
[root@zhuangdizhu2 init.d]#
```

Step4: start a container

```
[root@zhuangdizhu2 init.d]# docker run -i -t --name test ubuntu /bin/bash
```

```
Status: Downloaded newer image for ubuntu:latest
```

```
root@526e5d1ac302:/# apt-get install openssh-server
```

```
root@526e5d1ac302:/#apt-cache search "rdma"
```

```
Then install all of them
```

```
root@526e5d1ac302:/#apt-cache search "libib"
```

```
Then install those relevant to "infiniband"
```

```
root@526e5d1ac302:/#apt-cache search mlx
```

```
Then install all of them
```

```
root@526e5d1ac302:/#apt-cache search ibverbs-utils
```

```
root@526e5d1ac302:/#apt-cache search rdmacm-utils
```

```
root@526e5d1ac302:/#apt-cache search opensm
```

```
root@526e5d1ac302:/#apt-cache search libopensm-dev
```

```
Then install all of them
```

```
root@526e5d1ac302:~#exit
```

Step5: update and commit images

```
[root@zhuangdizhu2]# docker commit 526e5d1ac302 jude/ubuntu:jude1
```

```
[root@zhuangdizhu2]#docker ps -a
```

Step6: remove the older container

```
[root@zhuangdizhu2]#docker rm -f 526e5d1ac302
```

Step7: start a new container

```
[root@zhuangdizhu2]#docker run -i -t --privileged --name judy2 --  
device=/dev/infiniband/uverbs0 --device=/dev/infiniband/rdma_cm --  
device=/dev/infiniband/ucm0 --device=/dev/infiniband/umad0 jude/Ubuntu:dude1 /bin/bash
```

```
root@f32d28fb7b5b:/# ibv_rc_pingpong
```

```
local address: LID 0x0001, QPN 0x040225, PSN 0x0f9e19, GID ::
```

The container on the other host:

```
root@116c06041b87:/# ibv_rc_pingpong 172.16.1.3
```

```
local address: LID 0x0002, QPN 0x040229, PSN 0x04e46e, GID ::
```

```
Couldn't connect to 172.16.1.3:18515
```

(then got a port number: 18515)

exit, commit the image to jude/ubuntu:jude2, remove the stopped container, and start a new container with the following command:

```
docker run -i -t --privileged --name judy2 -p: 18515:18515 --device=/dev/infiniband/uverbs0 --  
device=/dev/infiniband/rdma_cm --device=/dev/infiniband/ucm0 --  
device=/dev/infiniband/umad0 jude/Ubuntu:dude1 /bin/bash
```

ibv_rc:

```
root@4cfe1cdc1d78:/# ibv_rc_pingpong 172.16.1.2
```

```
local address: LID 0x0001, QPN 0x040226, PSN 0x16f267, GID ::
```

```
remote address: LID 0x0002, QPN 0x04022b, PSN 0xea069e, GID ::
```

```
8192000 bytes in 0.01 seconds = 10911.75 Mbit/sec
```

```
1000 iters in 0.01 seconds = 6.01 usec/iter
```

rdma:

```
root@4cfe1cdc1d78:/# rdma_client -s 172.16.1.2
```

```
rdma_client: start
```

```
rdma_client: end 0
```

```
root@4cfe1cdc1d78:/#
```